

ABHINANDAN KUMAR

Data Engineer | SQL | Snowflake | dbt | Databricks

[Abhinandan Kumar](#) | +91-9817750262 | [linkedin.com/in/abhinandankumar](https://www.linkedin.com/in/abhinandankumar) | github.com/abhinandankumar

PROFESSIONAL SUMMARY

Data Engineer with 3 years of experience across SQL, ETL/ELT pipelines, cloud data platforms and data-driven applications. Hands-on with Snowflake and dbt for staged-to-mart transformations, incremental processing and analytical data modeling, and with PySpark/Databricks for IoT telemetry pipelines. Strong SQL foundation across Snowflake, SQL Server and PostgreSQL, with experience supporting production data workflows, performance investigation, reporting and application integration.

SKILLS & TECHNOLOGIES

Data Platforms & Databases: Snowflake, MS SQL Server, PostgreSQL, Databricks, Delta Lake

Snowflake & SQL: Snowflake SQL, Stages, COPY INTO, Streams, Tasks, Query Profiling, Virtual Warehouses, MERGE Patterns, T-SQL, CTEs, Window Functions, Stored Procedures, Views, Functions

Data Engineering: dbt, ETL/ELT, PySpark, Incremental Processing, Dimensional Modeling, Fact & Dimension Tables, Data Quality

Programming & Cloud: Python, SQL, Node.js, FastAPI, AWS S3, AWS Lambda, Kafka, RabbitMQ

Analytics & Tools: Power BI, Git, Azure DevOps, Databricks Unity Catalog, TOAD, SQL Developer

Dev Tools: Git, Databricks, Snowsight, TOAD, SQL Developer, Azure DevOps

WORK EXPERIENCE

Vantiva India | *Software Engineer* | *Data Engineer*

Jun 2024 – Present

Smart Spaces Snowflake · dbt · SQL Server · PySpark · Databricks · AWS · Power BI

- Built Snowflake ELT ingestion using stages, COPY INTO, Streams, and Tasks, enabled scheduled incremental loads from raw/staging tables to curated analytical datasets.
- Developed dbt SQL models across staging, intermediate and mart layers, using incremental materializations and surrogate-key joins to implement reusable transformations for downstream analytics.
- Developed analytical data models by defining source granularity and business keys, created organized datasets to help the business easily track key performance metrics and analyze operations.
- Debugged failed Snowflake loads and scheduled Tasks through load, task and query history; investigated schema and file-format errors, validated reruns, and used query profiles and warehouse metrics to tune recurring workloads.
- Developed SQL Server stored procedures, triggers and views, developed complex T-SQL for multi-table joins, CTE-based transformations and reporting workflows, while supporting historical-data archival and retention.
- Built ETL pipelines using PySpark and Databricks to process 10,000+ daily telemetry data points from IoT devices, automatically calculating Wi-Fi mesh uptime, replacing a manual, router-by-router process.
- Built an AWS Lambda-powered, on-demand reporting feature for the Smart Spaces web app, letting storage site managers generate Excel reports (last 24 hours or a custom date range) on router and IoT device performance across their sites, also used by the engineering team for monitoring and debugging across 50+ facility sites.
- Used these reports to debug connectivity and communication issues between LoRa and Zigbee-enabled sensors and routers such as routers going offline, devices failing to sync, or missing data packets, helping maintain a ~99% uptime SLA.
- Built real-time analytics dashboards in Power BI for sales teams, senior leadership, and clients to monitor device performance and usage statistics.

HomeSight Care Node.js · Kafka · In-memory Caching · RBAC/ABAC

- Migrated core APIs from direct database reads to a Kafka-backed in-memory caching layer with indexed lookups and database fallback, reducing high-traffic response times from ~7 seconds to 40–50 ms.
- Integrated RBAC/ABAC permission models into the migrated API layer, covering all user-module endpoints (personal data, hub information, and related APIs) to enforce consistent, role-based access control across the user base.
- Resolved 30+ bugs and delivered 10+ frontend features, improving application stability and user experience.

Wow Labz | *Backend Engineer Intern*

Apr 2024 – Jun 2024

Cab Booking Application Python · FastAPI

- Built a FastAPI-based reporting service used daily by admins and 10-12 drivers to track driver performance, revenue, and bookings - giving admins fleet-wide visibility and drivers self-service income tracking.
- Used Kafka-based messaging to move video data through a 7-phase pipeline - chunking video and generating transcripts, subtitles, and dubbed audio across multiple languages.

Yamaha Motor Solutions India | *Graduate Engineer Trainee* | *Full Stack Developer*

Jul 2023 – Apr 2024

Peer Review Application Node.js · React · GraphQL · PostgreSQL · RabbitMQ

- Built a full-stack peer review application adopted by a 30-engineer team (later split into sub-teams), replacing manual review tracking with automated status visibility for engineers and managers, integrated RabbitMQ for real-time reviewer notifications.
- Developed GraphQL APIs using the Hasura GraphQL engine on PostgreSQL, authoring complex SQL functions, cursors, and dynamic DML queries to fetch and populate information on frontend forms.

CERTIFICATIONS

- Databricks Data Engineering Professional & Databricks Fundamentals Accreditation
- McKinsey Forward Learning Program – Problem Solving, Leadership & Communication

EDUCATION

B.Tech in Computer Science Engineering

Kurukshetra University | GPA: 8.02 / 10

2019 – 2023